



Coregenix - Article 1 - The Hidden Cost of Cloud: Why Data Egress Deserves More Attention

Cloud cost optimisation has traditionally focused on right-sizing virtual machines, shutting down idle resources, and selecting appropriate storage tiers. While these remain important, they often overlook one of the fastest-growing contributors to cloud spend, data egress.

Whenever data leaves a cloud provider, whether to another cloud, an on-premises environment, or the public internet, it can incur additional charges. Unlike fixed infrastructure costs, egress charges are usage-based, making them difficult to forecast and even harder to control.

According to Akave's 2026 Data Engineer's Guide, cloud providers can charge several times more to retrieve data than to store it, making data movement one of the largest hidden contributors to cloud expenditure for many organisations.

For businesses running analytics platforms, AI/ML workloads, backups, disaster recovery, and hybrid cloud environments, these charges can quickly rival or even exceed compute costs.

Where Cloud Costs Usually Escalate

Cloud overspending is rarely caused by a single issue. Instead, it's typically the result of several inefficiencies accumulating over time.

The most common cost drivers include:

- Idle or oversized compute resources
- Storage that remains in premium tiers longer than necessary
- Cross-region and cross-cloud data transfers
- Data egress charges
- Limited visibility into cloud usage across teams and business units

While organisations often monitor compute utilisation closely, data movement is frequently overlooked. As cloud adoption grows, data transfer increases alongside it, creating unpredictable monthly costs that are difficult to budget for.

Cloud Migration: The Best Time to Optimise

A cloud migration isn't simply an infrastructure project; it's one of the best opportunities to improve long-term cost efficiency.



Rather than migrating existing workloads as-is, organisations should evaluate how applications communicate, where data resides, and how frequently information moves between services.

A successful migration strategy should include:

- Mapping application and data flows not just workloads
- Identifying cross-cloud and cross-region dependencies
- Comparing total usage-based costs instead of headline compute pricing
- Implementing tagging, budgets, and cost monitoring from day one

Many organisations optimise performance during migration but postpone cost governance.

Unfortunately, expensive architectural decisions become much harder to reverse once workloads are in production.

Why No-Egress Self-Provisioning Platforms Are Gaining Attention

As cloud spending continues to rise, organisations are increasingly looking beyond traditional hyperscaler pricing models.

According to Gartner's Worldwide Public Cloud End-User Spending Forecast (2026), global cloud spending is expected to increase by more than 20% year over year, making predictable cloud economics a growing business priority.

One solution gaining momentum is the adoption of self-provisioning platforms with minimal or zero egress fees.

These platforms typically provide:

- Transparent usage-based billing
- Zero or flat-rate data egress
- Self-service infrastructure provisioning
- Standards-based APIs for easier migration and interoperability

The biggest advantage isn't simply reducing costs—it's improving cost predictability.

Research published by EgressCost.com comparing leading cloud providers demonstrates that organisations moving large volumes of data can save substantially by using platforms that eliminate or significantly reduce egress charges. This becomes especially valuable for businesses operating AI platforms, media applications, analytics environments, or hybrid cloud infrastructures where data movement is constant.

Where This Approach Delivers the Greatest Value



While no-egress platforms aren't intended to replace every cloud workload, they can be particularly effective for data-intensive environments such as:

- Object storage
- Backup and archival systems
- Disaster recovery environments
- AI and machine learning datasets
- Analytics platforms
- Multi-cloud and hybrid cloud architectures

For organisations already invested in hyperscaler-native services, a hybrid strategy often delivers the best balance of performance and cost efficiency—keeping compute where it performs best while optimising storage and data-heavy workloads using more predictable pricing models.

A Practical Starting Point

Reducing cloud costs doesn't require a complete platform migration.

A practical approach includes:

- Reviewing cloud spend across compute, storage, networking, and data transfer
- Identifying workloads generating the highest egress costs
- Piloting a zero-egress storage platform for backup or archival workloads
- Comparing actual invoices over multiple billing cycles
- Expanding adoption based on measurable savings

As highlighted in the nOps AWS Egress Cost Guide (2026), organisations should evaluate cloud costs using real operational usage patterns rather than relying solely on published pricing calculators.

Looking Ahead

Cloud optimisation is evolving.

It is no longer enough to simply reduce compute usage or purchase reserved instances. As organisations embrace AI, hybrid cloud, edge computing, and multi-cloud architectures, data movement is becoming one of the most important drivers of cloud economics.

Organisations that proactively understand and optimise these costs will be better positioned to improve financial predictability, support innovation, and build more resilient cloud environments.

The Bottom Line



Cloud cost optimisation is no longer just about running fewer virtual machines or purchasing discounted compute capacity. It's about designing architectures that move data efficiently and minimise unnecessary operational costs.

Evaluating platforms with transparent pricing and minimal or zero egress fees can help organisations improve cost control, simplify budgeting, and build cloud environments that are both scalable and financially sustainable.

At CoreGeniX, we help organisations design secure, scalable, and cost-efficient cloud environments, from cloud migration and infrastructure modernisation to cloud governance and optimisation. Whether you're planning your first migration or looking to optimise an existing cloud strategy, understanding the true cost of data movement is an essential step towards long-term cloud success.

Sources Referenced

- Gartner – *Worldwide Public Cloud End-User Spending Forecast, 2026*
- Akave – *What Are Egress Fees? The 2026 Data Engineer's Guide*
- EgressCost.com – *Cloud Provider Egress Pricing Comparison (2026)*
- nOps – *AWS Egress Costs in 2026: How to Reduce Them?*